

Overview / Introduction

◆ **Title:**

VOCAHIRE – Real-Time Voice-Based Interview Coach using Agentic AI

◆ **Objective:**

VOCAHIRE is an intelligent, real-time **voice-based interview simulator** designed to prepare users for technical and behavioral interviews. It uses **voice input, LLM-based question generation, tone analysis, and feedback systems** to coach candidates interactively.

◆ **Problem Statement:**

Many candidates struggle with:

- Public speaking or hesitation during interviews.
- Lack of real-time feedback.
- Limited access to personalized mock interviewers.

There is a need for an **automated, personalized interview coach** that can:

- Ask questions based on resumes.
 - Provide spoken and textual feedback.
 - Improve candidate confidence using tone/hesitation analysis.
-

◆ **Proposed Solution:**

- VOCAHIRE uses **voice input** for interview simulation.
- **LLMs** (like Claude, FLAN-T5, etc.) generate:
 - Resume-based questions
 - HR-style fallback questions
 - Smart feedback
- **Tone and pause detection** is integrated to assess hesitation.
- Results are saved in per-user interview logs for analysis.

◆ **Target Users:**

- Students preparing for placements/internships.
 - Job seekers.
 - Career counseling platforms.
-

◆ **Key Features:**

- Multi-user login/registration system.
- Resume upload and parsing.
- Voice-to-text (Deepgram / Whisper).
- Interview question generation (Resume + General).
- Real-time tone/hesitation scoring.
- Voice + Text feedback using LLM.
- Logs stored per user.




High-Level Architecture

Here's a **clear architecture overview** of your VOCAHIRE system, showing how all the components interact in real-time:

◆ **User Interface (Terminal)**

- User registers or logs in
 - Uploads resume
 - Answers interview questions via mic
 - Receives spoken feedback
-

◆ **2. Resume Processing Module**

- Reads uploaded **Resume.pdf**
- Uses **PyMuPDF (fitz)** to extract text
- Parses and categorizes into:
 -  Education
 -  Projects
 -  Skills
- Stores structured data as: **parsed_resume.json**

◆ **3. Real-time Interview Flow Engine**

- Drives the Q&A session dynamically
 - Tracks current question, stores all interactions
 - Pulls questions from:
 - Resume-based LLM prompts (Claude 3)
 - HR fallback prompts (Claude)
 - Maintains memory and history during interview
-

◆ **4. Question Generator (LLM Agent)**

- **Claude 3 Haiku** via **OpenRouter**
 - Resume-based technical questions
 - General HR-style fallback questions
 - Prompts are customized with parsed resume info
 - Also used for giving detailed **feedback**
-

◆ **Voice Engine**

Speech-to-Text:

- User speaks using microphone
- Processed via:
 - **Deepgram API** (primary)
 - or **OpenAI Whisper** (offline)

Tone Detection:

- VAD (Voice Activity Detection)
- Hesitation score computed with silence gaps

Text-to-Speech:

- AI reads each question + feedback out loud using **pyttsx3**
-

◆ **Logging & Memory Storage**

- Each user has their own folder:
`users/{username}/`
- Inside:
 - `Resume.pdf`

- `parsed_resume.json`
 - `logs/interview_timestamp.json`
- Logs include:
 - Question
 - Answer
 - Hesitation score
 - LLM feedback

◆ *Technologies Used*

Layer	Tools / APIs
LLM	Claude 3 (OpenRouter), Flan-T5
STT	Deepgram, Whisper
Resume Parsing	PyMuPDF
Voice	WebRTC VAD, pytsx3
Backend	Python (Modular Files)
Memory	JSON + Folder Storage

Component-Level Architecture

VocaHire is designed with a modular, agent-based architecture that promotes clarity, scalability, and easy debugging. Each component in the system is responsible for a specific function and communicates through well-defined interfaces.

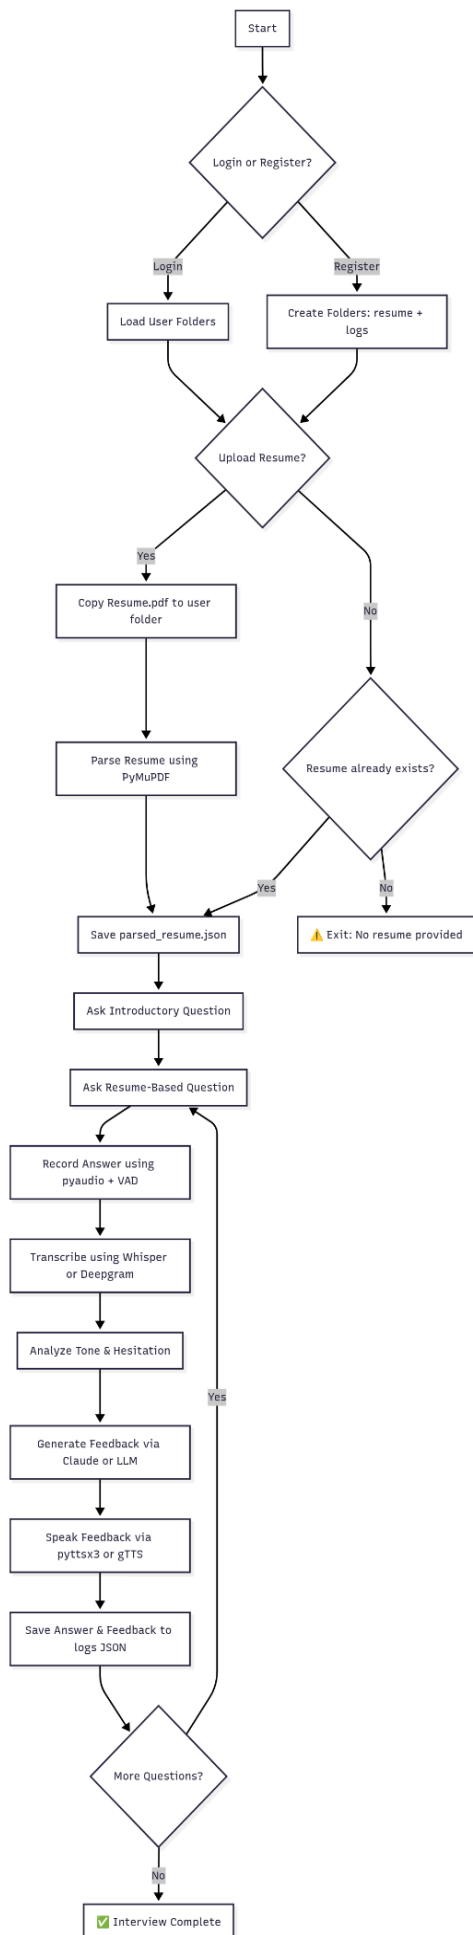
Core Components of VocaHire

Component	Role & Responsibility
1. User Manager	Manages login and registration. Creates per-user folders like <code>/users/{username}/resume</code> and <code>/users/{username}/logs</code> .
2. Resume Parser	Extracts structured data (education, skills, projects) from uploaded <code>Resume.pdf</code> using <code>PyMuPDF</code> . Saves output as <code>parsed_resume.json</code> .

3. Question Generator	Generates dynamic interview questions from three sources: <ul style="list-style-type: none"> • <code>resume_qa_generator.py</code> → Technical questions from resume • <code>general_qa_generator.py</code> → Introductory HR questions • <code>fallback_qa_generator.py</code> → Generic HR/behavioral questions
4. Interview Agent	Orchestrates the full interview flow: asking questions, collecting responses, tracking current state, and managing transitions.
5. Audio Input Handler	Records user's spoken response using <code>webrtcvad</code> . Applies voice activity detection for intelligent start/stop.
6. STT (Speech-to-Text)	Converts recorded <code>.wav</code> files to text using: <ul style="list-style-type: none"> • Local <code>Whisper</code> model or • <code>Deepgram API</code>
7. Tone Analyzer	Calculates hesitation/confidence score based on pauses and silence in the audio using a custom scoring algorithm.
8. Feedback Engine	Uses a Large Language Model (Claude 3 via OpenRouter) to evaluate answers and return helpful, structured feedback.
9. Logger	Stores the full interview session in a structured JSON format under <code>/users/{username}/logs/</code> , including question, answer, tone score, and feedback.

Interview Execution Pipeline

1. **User login/registration** creates or accesses personal workspace.
2. **Resume uploaded** and parsed into structured JSON.
3. **Interview starts:**
 - One **general intro** question first
 - Followed by **resume-aware** technical questions
 - Ends with **fallback HR/behavioral** questions
4. **Voice is recorded** and saved as `.wav`.
5. **STT engine transcribes** audio to text.
6. **Tone analysis computes** hesitation score.
7. **LLM generates feedback**.
8. **Session is saved** into a personal `.json` log file.



Component-Wise Explanation (Modular Design)

1. Login & User Manager

- **File:** `login.py` / main script
 - **Purpose:** Handles user registration, login, folder creation (resume, logs)
 - **Input:** Username
 - **Output:** Creates `users/<username>/resume`, `users/<username>/logs`
 - **Special Handling:** Prevents duplicates, prompts for resume upload
-

2. Resume Parser

- **File:** `resume/resume_parser.py`
 - **Library:** PyMuPDF (`fitz`)
 - **Purpose:** Extracts `education`, `skills`, `projects` from Resume.pdf
 - **Output:** `parsed_resume.json` saved to user folder
 - **Used In:** `resume_qa_generator.py`
-

3. Resume QA Generator

- **File:** `agents/resume_qa_generator.py`
 - **Model Used:** Claude via OpenRouter
 - **Purpose:** Generates technical interview questions from parsed resume
 - **Input:** `parsed_resume.json`
 - **Output:** 1 resume-based question per turn (customizable)
-

4. General QA Generator

- **File:** `agents/general_qa_generator.py`
- **Model:** Claude via OpenRouter
- **Purpose:** Generates intro/HR-style fallback questions like “Tell me about yourself”, “What are your strengths?”
- **Flow Control:** Comes **before** resume questions

5. Voice Capture & VAD

- **File:** `audio/step2_vad_listener.py`
 - **Library:** `webrtcvad`, `pyaudio`, `numpy`
 - **Purpose:** Records only voice segments using VAD
 - **Output:** `.wav` audio saved in user folder
-

6. Speech-to-Text

- **File:** `audio/stt_whisper_local.py` or `deepgram_transcriber.py`
 - **Backends Supported:** Whisper (local) / Deepgram (API)
 - **Purpose:** Converts voice to text for answering questions
-

7. Tone & Hesitation Analyzer

- **File:** `audio/tone_analysis.py`
 - **Purpose:** Computes hesitation % using silence gaps
 - **Logic:** High hesitation → Suggests improvement via feedback
-

8. Feedback Engine

- **File:** `agents/feedback_engine.py`
 - **LLM:** Claude (OpenRouter)
 - **Purpose:** Generates helpful, kind, actionable feedback
 - **Input:** question + answer + hesitation_score
 - **Output:** clear feedback string
-

9. TTS Speaker

- **File:** `audio/tts_speaker.py`
 - **Engines Supported:** `pyttsx3`, `gTTS`
 - **Purpose:** Speaks both questions and feedback aloud
 - **Customization:** Can set gender/voice rate
-

10. Interview Orchestration

- **File:** `audio/step4_voice_to_langgraph.py`
 - **Acts As:** Main loop controller
 - **Responsibilities:**
 - Asks questions in order (intro → resume → fallback)
 - Captures audio
 - Calls STT, analysis, feedback, and logging
-

11. Logging & Memory

- **Output:** `.json` logs stored in `users/<username>/logs/`
- **Fields Stored:** question, answer, hesitation %, feedback
- **Future Plan:** Can be queried by LangChain for memory-based interviews

Agent-Oriented Architecture (LangGraph-style)

Why Agent-Based?

Your system follows a **modular agent pattern**, where each module (question generation, STT, TTS, feedback, etc.) is an autonomous “agent” handling a single responsibility. This makes it easier to integrate with LangGraph in future or simulate its behavior manually.

Agents :

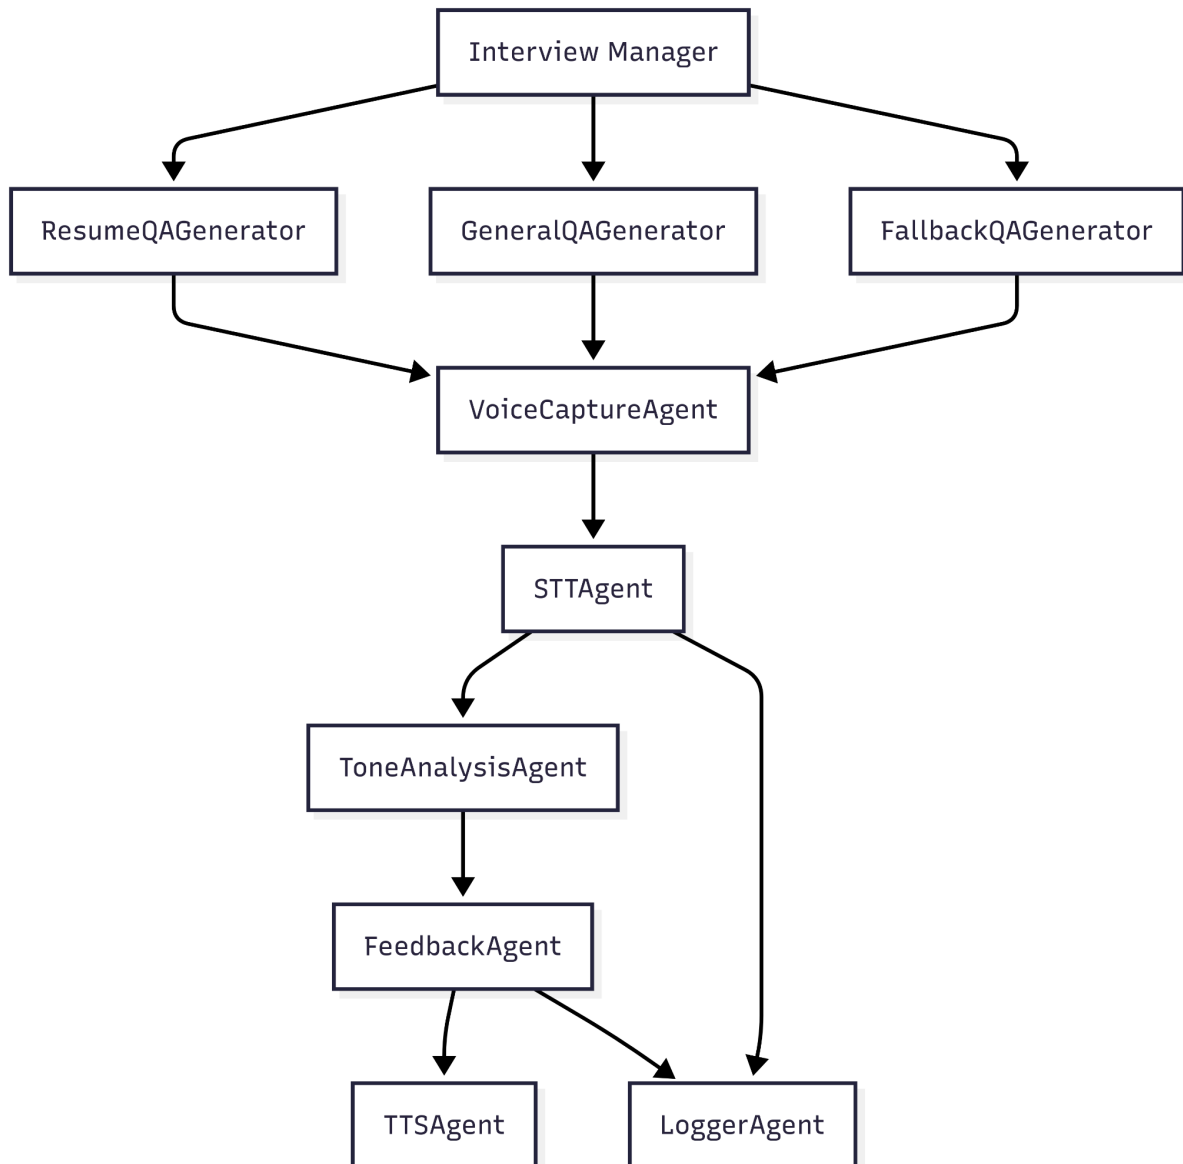
Agent Name	Responsibility	Input	Output
<code>ResumeParser Agent</code>	Extract resume data (skills, projects, etc.)	PDF resume	<code>parsed_resume.json</code>
<code>ResumeQA Generator</code>	Ask technical questions based on resume	Parsed resume	Resume-based question

VoiceCapture Agent	Record user audio using VAD	Microphone input	.wav file
STTAgent	Convert voice to text	Audio file	Answer text
ToneAnalysis Agent	Analyze pauses & compute hesitation score	Audio file	Hesitation score (%)
Feedback Agent	Generate feedback based on response tone & content	Answer + hesitation	Feedback string
TTSAgent	Speak question and feedback to the user	Text	Spoken audio output
Logger Agent	Store question, answer, score, and feedback	All outputs above	JSON log entry
Interview Manager	Coordinates flow: question → record → feedback	State tracker	Runs all agents in correct sequence

Future LangGraph Upgrade

You can replace `InterviewManager` with a real LangGraph flow using `langgraph.graph` and assign:

- `Stateful Mode` for resume memory
- `ConditionalEdge` to route based on score
- `LangChain` tools for each agent



Memory Strategy in VOCAHIRE

Memory Type	Method Used	Format	Where Stored
Interview Session	In-memory <code>InterviewState</code>	Python class	Tracks question index, state
Answers & Feedback	JSON file logs	<code>.json</code>	Stored per-user in <code>users/{name}/logs/</code>
Resume Parsing	Cached resume data	<code>.json</code>	<code>users/{name}/resume/parsed_resume.json</code>

Audio Files	Raw <code>.wav</code> recordings	<code>.wav</code>	<code>users/{name}/recordings/</code>
Multi-user Support	Folder structure	Folders	<code>users/{username}/...</code>

Core Components for Flow:

File	Purpose
<code>interview_agent.py</code>	Controls question flow
<code>resume_qa_generator.py</code>	Generates technical/resume questions
<code>fallback_qa_generator.py</code>	Generates HR-style questions
<code>submit_answer()</code>	Stores question-answer-feedback
<code>get_next_question()</code>	Selects next question logic

Components Used:

Type	Library	Description
Voice Activity Detection (VAD)	<code>webrtcvad</code>	Detects when the user starts/stops speaking
Audio Recording	<code>pyaudio</code> , <code>soundfile</code>	Captures voice and saves it as <code>.wav</code>
Speech-to-Text (STT)	<code>Deepgram API</code>	Transcribes spoken answers to text
Text-to-Speech (TTS)	<code>pyttsx3</code>	Speaks both questions and feedback










Files Involved:

File	Role
<code>step2_vad_listener.py</code>	Records voice based on VAD
<code>stt_whisper_local.py</code> (or <code>Deepgram</code>)	Converts <code>.wav</code> to text
<code>tone_analysis.py</code>	Calculates hesitation/pauses

tts_speaker.py

Converts text to speech using
pyttsx3

Sample Voice Flow:

-  AI: Tell me about yourself.
 -  Speaking: Tell me about yourself.
 -  Speak now...
 -  Start Recording
 -  Stop Recording
 -  Saved: recordings/recording_xxx.wav
 -  Transcribing...
 -  Your Answer: I'm a student at IIT Patna...
 -  Hesitation Score: 21.7%
-

Summary

VOCAHIRE is an AI-powered **real-time voice interview coach** designed to help users practice technical and behavioral interview questions in an interactive and personalized way. The system uses a combination of **local and API-based speech recognition, tone analysis, LLM-driven question generation, and smart feedback engines** to simulate the real interview experience.

As the developer, my goal was to create an intelligent, multi-user platform where candidates can:

- Speak naturally and be transcribed accurately
- Receive questions based on their **resume and HR scenarios**
- Get **real-time feedback** on tone, hesitation, and answer quality
- Maintain a record of their interviews and performance

This solution combines modular agents, natural language models (Claude via OpenRouter), resume parsing logic, voice activity detection, and a personalized feedback loop to deliver a full-stack interview simulation experience. The project was developed using **Python 3.11+, Deepgram/OpenAI Whisper, PyMuPDF, WebRTC VAD, and pyttsx3**, with flexible JSON storage and a folder-based multi-user system.

Key Achievements:

- Real-time voice conversation with VAD (webrtcvad)

- Accurate transcription using Whisper / Deepgram
- Resume parsing & dynamic question generation
- HR fallback & tone-aware feedback via Claude
- Multi-user support with session logging

Future Scope:

- Cloud-based deployment (e.g., using Streamlit, FastAPI)
 - Interview dashboard with analytics
 - Integration with ATS platforms or EdTech products
 - Support for group or panel interviews
-

Conclusion

VOCAHIRE is a modular, AI-driven voice interview platform designed to simulate real-time interview experiences. Through the integration of advanced speech-to-text (STT), large language models (LLMs), and emotion-aware feedback systems, it provides users with personalized coaching and actionable feedback.

By supporting resume parsing, memory, and multi-user capabilities, VOCAHIRE mimics the structure and flow of real interviews — making it a powerful tool for students and professionals preparing for technical roles.

This system is extensible, scalable, and can easily integrate future upgrades like better UI, database-backed memory, or multimodal feedback systems.